L-Università ta' Malta

Ensuring Correctness in Distributed Systems

An example report

by

Duncan Paul Attard

Department of Computer Science
Faculty of ICT

supervised by

ADRIAN FRANCALANZA
LUCA ACETO
ANNA INGÓLFSDÓTTIR

September 30, 2020

Faculty of ICT

Declaration

I, the undersigned, declare that the dissertation entitled:

Ensuring Correctness in Distributed Systems

submitted is my work, except where acknowledged and referenced.

Duncan Paul Attard

September 30, 2020

Abstract

A number of software systems today are built in terms of independently executing components that typically reside on different physical locations. While these software organisations offer a number of advantages, including the use of replication to improve robustness and quality of service, they are hard to design and implement. Ascertaining their correctness, therefore, becomes a chief concern. Traditional formal verification techniques, such as testing or model checking, tend to be applied with limited success in these scenarios due to a number of reasons. Runtime verification may be employed as a lightweight and dynamic alternative that complements the aforementioned verification approaches.

. . .

Contents

List of Figures

Chapter 1

Introduction

Numerous software systems are nowadays architected in terms of *asynchronous components* [2, 7, 1] that execute independently to one another without recourse to a global clock or shared state. Instead, components interact together via well-defined interfaces and non-blocking messaging [6] to create dynamic and loosely-coupled software organisations. Such architectures facilitate incremental updates, tolerate independent component failures and permit the various units of execution to be *distributed* across different locations [9, 3]. Despite their advantages, these systems are notoriously hard to design, and even harder to program and get right, and ensuring their correctness in terms of their expected behaviour becomes paramount.

1.1 Distributed Runtime Verification

While distributed systems inherit the characteristics inherent to asynchronous settings, their execution is further complicated due to physical constraints, such as the lack of a global clock and possibility of independent failures [5, 3]. In the local asynchronous case, traditional pre-deployment verification techniques such as model checking and testing [8, 11] often *scale poorly* because the set of execution paths considered is invariably dwarfed by the vast number of possible execution paths of the system. In a distributed scenario, use of these verification approaches is often problematic, if not *impractical*, due to the aforementioned complications.

Runtime Verification (RV) [10, 4] is complementary approach that evades some of the limitations of pre-deployment techniques by deferring the analysis until runtime. It employs *monitors* to *incrementally* analyse the system's behaviour (exhibited as a sequence of *trace* events) up to the current execution point, to determine whether a correctness specification under investigation is satisfied or violated.

. . .

Example 1.1. Consider the system SYS composed of processes P and Q, given in terms of the *labelled transition systems* in figure 1.1. P can initially fork (denoted by action f) a process, after which it either sends (s) a message and exits (e), or exits (e) immediately; Q is a simpler process that performs a single receive (r) and exits (e). A possible correctness property states that *"SYS does not fork processes at startup"*.

When SYS exhibits the witness trace *f.r.e.e*, the monitor can detect a violation of this property. For a different execution interleaving, *e.g. r.e.f.e* (where *f* is not the first event), the typical RV analysis would be unable to detect the fact that SYS is capable of performing f. Recouping this *lack of precision* is possible, but this would require the specification to consider *all* the possible trace event permutations that the composition of P and Q may exhibit (figure 1.1b). One easily observes that adding new components to SYS aggravates the specification task to the point where it becomes unwieldy and error-prone. Reformulating the original property to consider P in isolation, *i.e., "P does not fork processes at startup"*, eliminates the need to account for the behaviour of other nonrelevant components. ∎
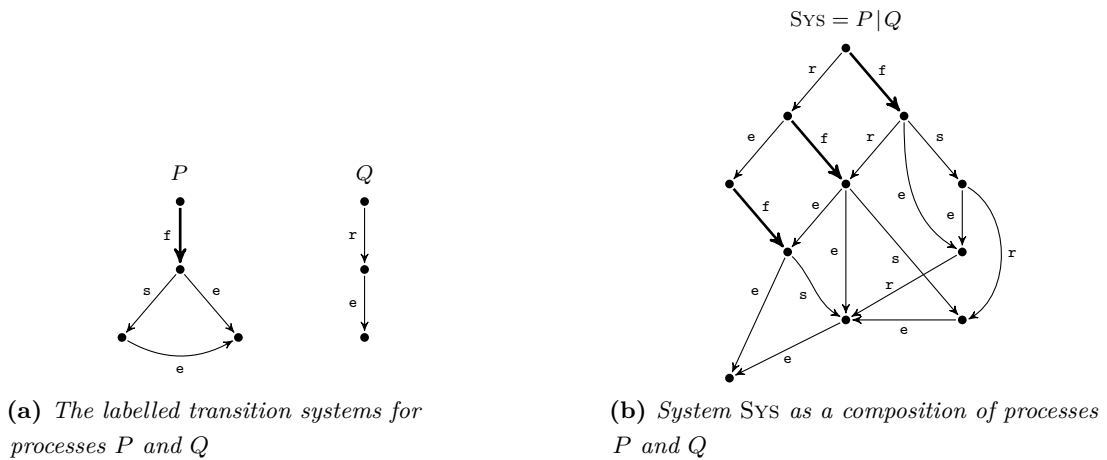
. . .

(a) *The labelled transition systems for processes P and Q*

(b) *System SYS as a composition of processes P and Q*

Figure 1.1. *Local and global states of a component-based system*

Chapter 2

Background

This chapter presents an overview of the background literature and state-of-the-art in the field of RV for distributed systems. We first give a brief synopsis of RV, and define the terms introduced in the preceding chapter. This is followed by an account of distributed systems that acquaints readers with the concepts used throughout this report. Finally, a series of recent works is discussed, comparing and contrasting similarities and differences between them. This exposition should reassert the current limitations in the area, and in doing so, underscore the novelty of our research contributions proposed in chapter 1.

. . .

Bibliography

[1] G. Agha, I. A. Mason, S. F. Smith, and C. L. Talcott. A Foundation for Actor Computation. *JFP*, 7(1):1–72, 1997.

[2] D. Chappell. *Enterprise Service Bus: Theory in Practice.* O'Reilly Media, 2004.

[3] J. Dollimore, T. Kindberg, and G. Coulouris. *Distributed Systems: Concepts and Design.* Addison Wesley, 2005.

[4] Y. Falcone, J. Fernandez, and L. Mounier. What can you verify and enforce at runtime? *STTT*, 14(3):349–382, 2012.

[5] S. Ghosh. *Distributed Systems: An Algorithmic Approach.* Chapman and Hall/CRC, 2014.

[6] G. Hohpe and B. Woolf. *Enterprise Integration Patterns: Designing, Building, and Deploying Messaging Solutions.* Addison-Wesley Professional, 2003.

[7] N. M. Josuttis. *SOA in Practice: The Art of Distributed System Design: Theory in Practice.* O'Reilly Media, 2007.

[8] E. M. C. Jr., O. Grumberg, and D. A. Peled. *Model Checking.* The MIT Press, 1999.

[9] G. V. K. *Elements of Distributed Computing.* Wiley India, 2014.

[10] M. Leucker and C. Schallhart. A brief account of runtime verification. *JLAP*, 78(5):293–303, 2009.

[11] G. J. Myers, C. Sandler, and T. Badgett. *The Art of Software Testing.* Wiley, 2011.

Appendix A

Trace Partitioning and Local Monitoring for Asynchronous Components

The paper entitled *"Trace Partitioning and Local Monitoring for Asynchronous Components"* was published in SEFM 2017 under Springer LNCS.

...